


1 Overview

You are working as a Developer for LetsManage Inc. The development team has decided to implement a Proof of Concept for a new Memory Management Unit which needs to be developed for a new Operating System version. You must implement the 4 APIs described below. Your code will be tested against a series of calls to these APIs and the output. NOTE: insert(),remove() are invoked by the OS itself while write(),read() are user calls.

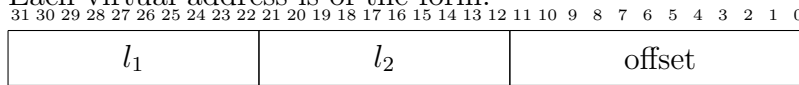
You are going to implement a 2-level page table for a 32-bit architecture.

- Each page is 4096 bytes,
- Each *Page Table Entry (PTE)* is 32 bits



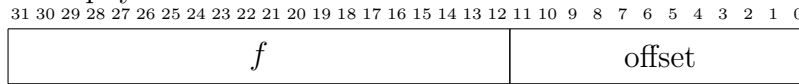
- The top 20 bits are the frame **f** and
 - The bottom 12 bits are the permissions (only the low order 3 bits are used)
 - The permissions are (by bit position):
 - 2:** R (read),
 - 1:** W (write), and
 - 0:** P (present)
 - The contents of the PTE are a 32 bit unsigned int, but we shall write it as a pair (*f, perm*) where *f* is the frame and *perm* is the permissions. The permissions will be written here using **RWP** notations.
- For example as **R-P** which means permissions R and P, but not W. Numerically, bit 2 and 0 are set, and thus the value is 5. Thus (17,5) has the numerical value 17*4096+5.

- Each virtual address is of the form:



- $l_1||l_2$ is the page number. Page numbers are 20-bit values.

- Each physical address is of the form:



- f is the frame. Frames are 20-bit values, although physical memory is smaller than this.

2 Interfaces and data structures

You should have the definitions as shown in Figure 1. Please read them carefully, they are critical for understanding the statements below.

- `pageTable` contains page table (enough to represent a single process). The `pageTable` is physically addressed.
- `userPages` contain the actual values which are read/written from user space. The `userPages` are physically addressed.
- `insert_pte` inserts the page table entry (mapping) in the page table from page `p` to frame `f`.
 - The 1st level page table entry will have all three permissions set, ie RWP.
 - The 2nd level page table entry set by the insert will have permissions `read_vaddr` and `write_vaddr` as given in the arguments and will also sent P.

It returns 1 if successful. For errors, see the below error list.

- `remove_pte` removes the page table entries for page `p`. The 2nd level page table entry is set to 0 i.e. (0, ---). It returns 1 if successful. For errors, see the below error list.
- `read_vaddr` reads a character at physical address `paddr` in `userPages`, after using the page table to map `vaddr` to `paddr`. If any page table entry used in the translation does not have both R and P bits set, the translation fails. It returns 1 if successful. For errors, see the below error list. On success, it returns the character read as `*c`.
- `write_vaddr` writes a character `c` n at physical address `paddr` in `userPages`, after using the page table to map `vaddr` to `paddr`. If any page table entry used in the translation does not have W and P bit sets, the translation fails. It returns 1 if successful. For errors, see the below error list.

3 Errors

The error codes are as follows:

- -1: missing read permission
- -2: missing write permission
- -3: not present
- -4: 2nd level page table entry already defined (when doing an insert)

```

#define PAGE_SIZE_IN_BYTES 4096
#define PAGE_SIZE_IN_UINTS 1024
#define PAGE_TABLE_PAGES 1025
#define USER_PAGE_PAGES 1000

typedef unsigned int uint32;

/* a single page of the page table holds
   1024 page table entries, each a uint
   */
typedef uint32 PtePage[1024];

/* the page table consists of 1025 pages:

   . pageTable[0] holds the level1 page table page

   . pageTable[1]...pageTable[1024]
     holds the 1024 level2 page table pages
   */
PtePage pageTable[PAGE_TABLE_PAGES];

/* holds the user pages, byte addressable, where
   userPages[0].. userPages[4095] holds the first userPage
   */
char userPages[PAGE_SIZE_IN_BYTES*USER_PAGE_PAGES];

int insert_pte(uint32 p, uint32 f, uint32 read, uint32 write);
int remove_pte(uint32 p);
int read_vaddr(uint32 vaddr, char *c);
int write_vaddr(uint32 vaddr, char c);

```

Figure 1: Declarations

4 Example

4.1 Page table walk

Consider a page table walk for the above virtual address for a **read**.

address	Page Table Entry	
	frame	permission check
pageTable[0][l_1]	f_1	R and P
pageTable[f_1][l_2]	f_2	R and P
userPage[$f_2 \ll 12 + \text{offset}$]		

4.2 Sequence of operations

Next, we consider a sequence of calls, including the most critical procedure `insert()`.

The key to this assignment is the `insert()` procedure call. Note that the page number `p` has the range 0x00000 to 0xFFFFF, and that both `pageTable` and `userPages` are accessed by physical addresses.

The `pageTable` is split into one 1st level page at `pageTable[0]` and 1024 2nd level pages, where 2nd level page i is at address `pageTable[i+1]`. Thus there are 1025 page frames for the `pageTable`.

The `userPages` is byte addressible and has a 1000 frames, indexed in the range 0 to 999; this corresponds to physical addresses 0 ... 1000*4096-1. Obviously you should not map to a physical address outside this range.

i	call	effect	return
1	<code>insert_pte(23, 5, 1, 0);</code>	<code>pageTable[0][0] = (1, RWP);</code> <code>pageTable[1][23] = (5, R-P)</code>	1
2	<code>insert_pte(1023, 999, 1, 1)</code>	<code>pageTable[0][0] = (1, RWP);</code> <code>pageTable[1][1023] = (999, RWP)</code>	1
3	<code>insert_pte(4097, 29, 0, 1)</code>	<code>pageTable[0][4] = (5, RWP);</code> <code>pageTable[5][1] = (29, -WP)</code>	1
4	<code>write_vaddr(1023*4096+131, 'x')</code>	<code>userPages[999*4096+131] = 'x'</code>	1
5	<code>read_vaddr(1023*4096+131, cP)</code>		1, *cP='x'
6	<code>write_vaddr(23*4096, cP)</code>		-2

1. This will map page 23 ($l_1 = 0, l_2 = 23$) to frame 5 with read permission
2. This will map page 1023 ($l_1 = 0, l_2 = 1023$) to frame 999 with read and write permissions
3. This will map page 4097 ($l_1 = 4, l_2 = 1$) to frame 29 with write permission
4. This writes an 'x' at physical address $999*4096 + 131$
5. This read an 'x' (*cPtr) at physical address $999*4096 + 131$
6. This will fail because of a lack of write permissions

